

A Comprehensive Analysis of Life Expectancy Determinants: Data Cleaning, Model Specification, and Diagnostic Testing

Bhupathi Nikhil Agnihotri (230123074)	Ryan Aggarwal (230101117)	Nikhil Redkar (230123073)	Vipul Sharma (230103117)
Kunj Aggarwal (230103056)	Devansh Jangid (230102027)	Toshit Jain (230101106)	Dhruv Kondapi (230122024)

1 Introduction

Life expectancy is a crucial measure of public health and a key indicator of socioeconomic development. Understanding the determinants of life expectancy provides valuable insights into public policy and investment priorities. This report examines the factors affecting life expectancy using a dataset that includes economic, health, and environmental variables. Our study begins with a thorough data cleaning and normalization process, followed by a stepwise regression analysis. In addition, we perform various diagnostic tests to validate the model assumptions, ensuring that the results are robust and reliable.

1.1 Context and Background

Life expectancy is influenced by a myriad of factors:

- **Economic conditions:** Higher GDP per capita generally leads to better nutrition, improved healthcare services, and enhanced living standards.
- **Health indicators:** Variables such as infant mortality and health expenditure provide direct measures of a country's healthcare quality.
- **Environmental factors:** Air quality (measured by PM2.5) and access to clean water can impact public health.

Prior research has shown that improvements in these areas can lead to increased life expectancy. However, the simultaneous inclusion of multiple variables can lead to issues such as multicollinearity, heteroskedasticity, and influential observations that may distort the results. Our analysis carefully addresses these issues to arrive at a parsimonious and robust model.

1.2 Objectives

The main objectives of this report are:

- a. **Data Preparation:** To clean, normalize, and prepare the dataset for regression analysis by handling missing values, eliminating redundant variables, and scaling the data.
- b. **Model Estimation and Variable Selection:** To build an OLS regression model that explains life expectancy using a set of selected predictors and to justify the inclusion or elimination of certain variables.
- c. **Diagnostic Testing:** To perform tests for multicollinearity, heteroskedasticity, and normality of residuals, along with an assessment of influential observations.
- d. **Interpretation:** To discuss the implications of the model coefficients and diagnostic tests, including the effects of outlier removal, and to identify the most optimal model parameters.

2 Data and Methodology

2.1 Data Overview

The dataset used in this study is sourced from international databases such as the World Bank and national health surveys. It includes:

- **Economic Indicators:** GDP per capita (`gdp_pc`) is included as a proxy for economic prosperity.
- **Health Metrics:** Variables like infant mortality (`infant_mortality`) and health expenditure (`health_expenditure`) directly relate to healthcare quality.
- **Environmental and Social Indicators:** Access to water (`water_access`), air quality (PM2.5, denoted as `pm25`), the number of physicians, and employment rate provide additional context.
- **Dependent Variable:** Life expectancy (`life_expectancy`), defined as the average number of years a newborn is expected to live given current mortality conditions.

2.2 Data Cleaning and Normalization

Data cleaning involved the following key steps:

- **Importing and Renaming:** The raw data was imported from a CSV file. Variable names were shortened for ease of interpretation (e.g., `gdp_pc` instead of GDP per capita).
- **Handling Missing Values:** Placeholder values were replaced with Stata’s missing symbol. Variables originally stored as strings were converted to numeric formats.
- **Dropping Irrelevant Variables:** Variables such as employment rate, water access, and physicians were initially included but later dropped after analysis showed they were not statistically significant or introduced collinearity.

Normalization was applied to ensure that all predictors (except life expectancy) were on a similar scale. A min-max transformation was used:

$$X_{\text{norm}} = \frac{X}{\max(X)} \times \text{scale},$$

where *scale* is the scaling factor of the variable (10 in our case). This approach maintains the distributional shape while ensuring comparability among predictors.

2.3 Variable Selection and Correlation Analysis

Initial regression models included all candidate predictors:

$$\begin{aligned} \text{life_expectancy} = & \beta_0 + \beta_1 \cdot \text{gdp_pc} + \beta_2 \cdot \text{infant_mortality} + \beta_3 \cdot \text{health_expenditure} \\ & + \beta_4 \cdot \text{water_access} + \beta_5 \cdot \text{pm25} + \beta_6 \cdot \text{physicians} + \beta_7 \cdot \text{employment_rate} + \epsilon. \end{aligned}$$

After examining the statistical significance of the coefficients, variables such as `water_access`, `physicians`, and `employment_rate` were dropped. The rationale for elimination was:

- **Lack of Statistical Significance:** Employment rate showed a p-value of 0.654, indicating that it did not contribute significantly to explaining life expectancy.
- **Collinearity Concerns:** Although the Variance Inflation Factor (VIF) values were low overall, the combined presence of these variables did not improve the model’s explanatory power.

Thus, the model was refined to:

$$\text{life_expectancy} = \beta_0 + \beta_1 \cdot \text{gdp_pc} + \beta_2 \cdot \text{infant_mortality} + \beta_3 \cdot \text{health_expenditure} + \beta_4 \cdot \text{pm25} + \epsilon.$$

2.4 Outlier Identification and Removal

Influential observations were identified using Cook's distance. The formula for Cook's distance for observation i is given by:

$$D_i = \frac{\sum_{j=1}^n (\hat{y}_j - \hat{y}_{j(i)})^2}{p \cdot MSE},$$

where $\hat{y}_j$ is the predicted value using the full dataset, $\hat{y}_{j(i)}$ is the predicted value when observation i is omitted, p is the number of predictors, and MSE is the mean squared error. A threshold of 0.001 was applied, and observations exceeding this value were removed. Outlier removal resulted in a more homogeneous error distribution and improved model fit.

2.5 Model Summary Statistics and Diagnostic Tests

2.5.1 Coefficient of Determination (R^2)

The R^2 statistic measures the proportion of the variance in the dependent variable explained by the independent variables:

$$R^2 = 1 - \frac{SS_{\text{Residual}}}{SS_{\text{Total}}},$$

where SS_{Residual} is the sum of squared residuals and SS_{Total} is the total sum of squares.

2.5.2 Adjusted R^2

Adjusted R^2 accounts for the number of predictors and adjusts for the degrees of freedom:

$$\bar{R}^2 = 1 - \left(\frac{SS_{\text{Residual}}/(n - p - 1)}{SS_{\text{Total}}/(n - 1)} \right),$$

where n is the number of observations and p is the number of predictors.

2.5.3 F-test for Overall Significance

The F-test assesses the overall significance of the regression model:

$$F = \frac{(SS_{\text{Model}}/p)}{(SS_{\text{Residual}}/(n - p - 1))}.$$

A large F-value and a corresponding small p-value indicate that the model provides a significantly better fit than a model with no predictors.

2.5.4 Variance Inflation Factor (VIF)

VIF quantifies multicollinearity:

$$\text{VIF}_j = \frac{1}{1 - R_j^2},$$

where R_j^2 is the R^2 from regressing predictor j on all other predictors. A VIF below 5 (and in our case, below 2) is generally considered acceptable.

2.5.5 Breusch-Pagan/Cook-Weisberg Test for Heteroskedasticity

The Breusch-Pagan test is used to test the null hypothesis of constant variance (homoscedasticity):

$$\chi^2 = \frac{n \cdot R_{\text{aux}}^2}{2},$$

where R_{aux}^2 is the R^2 from an auxiliary regression of the squared residuals on the fitted values. A high p-value suggests that we cannot reject the null hypothesis of homoscedasticity.

2.5.6 Normality of Residuals (Q-Norm Test)

While there is no single formal formula for the Q-norm test, a quantile-quantile (Q-Q) plot is used to visually assess if the residuals follow a normal distribution. Deviations from the straight-line pattern indicate departures from normality.

3 Results and Discussion

3.1 Regression Outputs

Three regression models were estimated:

Model 1: Full Model (7 Predictors)

```
-----  
Regression Results: life_expectancy ~ gdp_pc + infant_mortality + health_expenditure  
                    + water_access + pm25 + physicians + employment_rate  
-----
```

Number of obs = 2133

F(7, 2125) = 1304.72, Prob > F = 0.0000

R-squared = 0.8112, Adj R-squared = 0.8106, Root MSE = 2.733

```
-----  
Variable          | Coef.      | Std. Err.   | t      | P>|t|      | [95% Conf. Interval]  
gdp_pc            | 0.9325091  | 0.0578491   | 16.12  | 0.000      | 0.8190624 1.045956  
infant_mortality  | -2.687371  | 0.0466072   | -57.66 | 0.000      | -2.778772 -2.595971  
-----
```

health_expenditure	1.231411	0.073379	16.78	0.000	1.087508	1.375313
water_access	0.3424513	0.0968956	3.53	0.000	0.1524312	0.5324714
pm25	-0.1497846	0.0562911	-2.66	0.008	-0.260176	-0.0393931
physicians	-0.2099529	0.0535144	-3.92	0.000	-0.3148989	-0.1050068
employment_rate	0.0299884	0.0668554	0.45	0.654	-0.1011205	0.1610973
_cons	69.52659	1.148211	60.55	0.000	67.27486	71.77833

Model 2: Reduced Model (After Dropping Non-significant Variables)

Regression Results: life_expectancy ~ gdp_pc + infant_mortality + health_expenditure + pm25

Number of obs = 2133

F(4, 2128) = 2250.61, Prob > F = 0.0000

R-squared = 0.8088, Adj R-squared = 0.8085, Root MSE = 2.7486

Variable	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]
gdp_pc	0.9348535	0.0570532	16.39	0.000	0.8229676 1.046739
infant_mortality	-2.663443	0.0447784	-59.48	0.000	-2.751257 -2.575629
health_expenditure	1.20086	0.0725143	16.56	0.000	1.058654 1.343066
pm25	-0.1685564	0.0494134	-3.41	0.001	-0.26546 -0.0716529
_cons	72.48601	0.3113791	232.79	0.000	71.87537 73.09665

Model 3: Final Model (After Removing Influential Outliers)

Final Regression Results: life_expectancy ~ gdp_pc + infant_mortality + health_expenditure + pm25

Number of obs = 1931

F(4, 1926) = 2539.82, Prob > F = 0.0000

R-squared = 0.8406, Adj R-squared = 0.8403, Root MSE = 2.2069

Variable	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]
gdp_pc	1.429883	0.0681422	20.98	0.000	1.296243 1.563524
infant_mortality	-2.5769	0.0457536	-56.32	0.000	-2.666632 -2.487168
health_expenditure	1.429138	0.0740367	19.30	0.000	1.283937 1.574338
pm25	-0.2105475	0.0505806	-4.16	0.000	-0.309746 -0.111349
_cons	71.39453	0.2977408	239.79	0.000	70.8106 71.97846

3.2 Interpretation of Results and Variable Selection

- **Inclusion of Variables:** The final model retains gdp_pc, infant_mortality, health_expenditure, and pm25. These variables are theoretically justified:

- **GDP per Capita:** Represents economic prosperity and correlates positively with improved healthcare and living standards.
 - **Infant Mortality:** Acts as a critical health indicator; higher infant mortality implies poorer overall healthcare and a shorter life expectancy.
 - **Health Expenditure:** Reflects investment in the healthcare system; higher spending is linked to better health outcomes.
 - **PM2.5:** Although its effect is negative, it is statistically significant in the final model, suggesting that air quality does have an impact after controlling for other factors.
- **Elimination of Variables:** Variables such as `water_access`, `physicians`, and `employment_rate` were dropped due to their lack of statistical significance or their introduction of collinearity without substantial improvement in model fit.

3.3 Impact of Outlier Removal

The removal of influential observations using Cook’s distance (threshold $D_i \geq 0.001$) led to:

- A decrease in the Root Mean Squared Error (from 2.7486 to 2.2069), indicating improved model precision.
- An increase in the R^2 from 0.8088 to 0.8406, demonstrating that the final model explains a larger proportion of the variation in life expectancy.
- Enhanced stability of coefficient estimates and a more reliable test for heteroskedasticity.

3.4 Diagnostic Test Summary

1. Multicollinearity:

- VIFs for the final model were as follows: `gdp_pc` (1.69), `pm25` (1.66), `infant_mortality` (1.55), and `health_expenditure` (1.30). All values are well below 5, indicating that multicollinearity is not a concern.

2. Heteroskedasticity:

- The Breusch-Pagan/Cook-Weisberg test yielded a chi-square value of 1.24 with a p-value of 0.2649 in the final model, suggesting constant error variance.

3. Normality of Residuals:

- Q-Q plots (generated by the `qnorm` command) showed that the residuals follow a straight line, implying approximate normality despite slight departures detected by formal tests.

3.5 Summary of the Most Optimal Model

The final, most optimal model is defined by the regression equation:

$$\text{life_expectancy} = \beta_0 + \beta_1 \cdot \text{gdp_pc} + \beta_2 \cdot \text{infant_mortality} + \beta_3 \cdot \text{health_expenditure} + \beta_4 \cdot \text{pm25} + \epsilon.$$

Key parameters include:

- $\beta_1 \approx 1.4299$: Suggesting that, all else equal, an increase in normalized GDP per capita is associated with a significant increase in life expectancy.
- $\beta_2 \approx -2.5769$: A one-unit increase in infant mortality is associated with a significant reduction in life expectancy.
- $\beta_3 \approx 1.4291$: A one-unit increase in health expenditure corresponds to a significant increase in life expectancy.
- $\beta_4 \approx -0.2105$: Even after controlling for economic and health factors, higher PM2.5 levels are linked to reduced life expectancy.

These coefficients, together with the improved R^2 and lower Root MSE, indicate that the final model is both statistically and theoretically significant.

4 Conclusion

4.1 Summary of Findings

This analysis provides a comprehensive understanding of the determinants of life expectancy. The final regression model, after careful variable selection and the removal of influential outliers, reveals that:

- **Economic prosperity (GDP per capita)** plays a crucial role in enhancing life expectancy.
- **Infant mortality** has a strong negative impact, underscoring the importance of maternal and child health interventions.

- **Health expenditure** positively affects life expectancy, supporting increased investment in healthcare.
- **Air pollution (PM2.5)**, although less influential compared to the other variables, still has a significant negative effect.

4.2 Implications for Policy and Business

Policy Implications:

- **Economic Policies:** Promoting economic growth and higher GDP per capita can indirectly improve health outcomes and extend life expectancy.
- **Healthcare Investment:** Increased spending in health infrastructure and services can directly improve public health and reduce mortality rates.
- **Child Health Initiatives:** Programs aiming to reduce infant mortality are critical and can have large positive effects on the overall life expectancy.

Business Implications:

- **Healthcare Sector:** Investors and businesses in the healthcare industry should note the positive impact of health expenditure, as it suggests a favorable environment for expansion.
- **Economic Growth and Productivity:** A healthier population, fostered by improved economic conditions, creates a more productive workforce, further aiding business growth.

4.3 Future Research Directions

Future studies might explore:

- **Inclusion of Additional Variables:** Factors such as education, environmental regulations, and lifestyle indicators may further explain variations in life expectancy.
- **Non-linear Relationships:** Testing for non-linear effects and interactions between predictors could yield deeper insights.
- **Advanced Methodologies:** Approaches such as robust regression or panel data analysis might help in further addressing endogeneity or time-specific effects.

4.4 Final Remarks

By carefully cleaning the data, normalizing variables, and addressing issues of multicollinearity, heteroskedasticity, and influential observations, this study has produced a robust model

for understanding life expectancy. The final model not only has strong statistical support—evidenced by high R^2 and adjusted R^2 values and significant F-tests—but it also aligns with theoretical expectations regarding the influence of economic and health-related factors.

In summary, the rigorous methodology employed in this analysis provides compelling evidence that policies aimed at boosting economic performance, reducing infant mortality, and increasing healthcare expenditure are likely to have a substantial positive impact on life expectancy. These findings are essential for guiding future policy decisions and strategic business investments in the public health domain.